

Here is a screenshot from tensorboard:

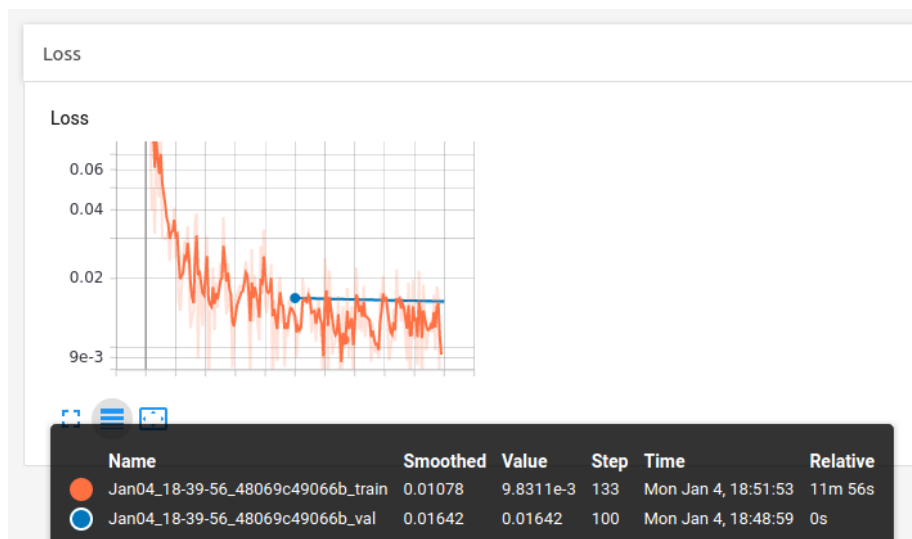




Figure 1: tensorboard screenshot

It shows the training loss is orange and the validation loss in blue. The model was trained for only two epochs; that's why there are only two validation loss points.

The following images of predictions during training were also obtained from tensorboard output:


Early stage of training (image, then ground truth, then prediction):







A middle stage of training (image, then ground truth, then prediction):





Towards the end of training (image, then ground truth, then prediction):





Here is a screenshot of slicer, while playing with an inference after the model was trained:

The 3D volume shown in the screenshot was predicted by the model. The 2D views show both the prediction and the ground truth segmentations overlayed.

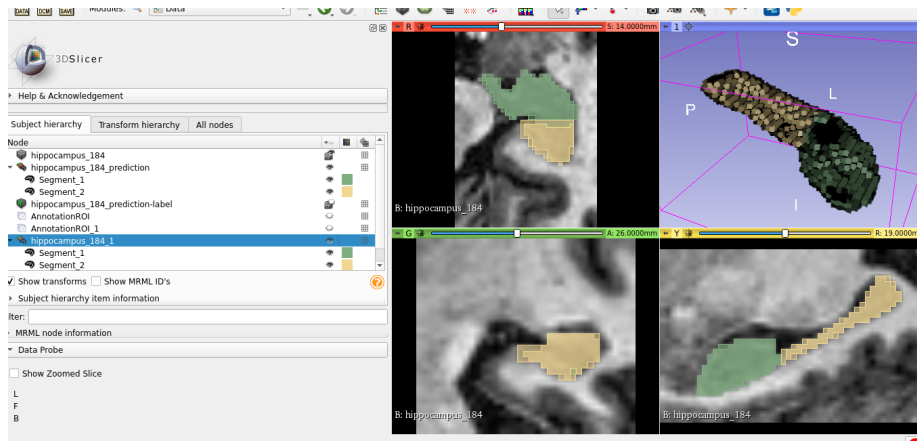


Figure 2: slicer screenshot